

Experiment: Kind tcpdump image load

Create Cluster

For MacOS and Windows

Launch a terminal in your Projects folder

Create a kind cluster with a [here document](#) for the cluster config.

```
$ kind create cluster --name tcpdump-cluster --image kindest/node:v1.21.12
```

```
👉 Preparing nodes 🚀 Creating cluster "tcpdump-cluster" ...
👉 Ensuring node image (kindest/node:v1.21.1) 🚀 ...
👉 Ensuring node image (kindest/node:v1.21.1) 🚀 ...
👉 Preparing nodes 🚀 ...
👉 Preparing nodes 🚀
👉 Writing configuration 🚀 ...
👉 Writing configuration 🚀
👉 Starting control-plane 🚀 ...
👉 Starting control-plane 🚀
👉 Installing CNI 🚀 ...
👉 Installing CNI 🚀
👉 Installing StorageClass 🚀 ...
👉 Installing StorageClass 🚀
Set kubectl context to "kind-tcpdump-cluster"
You can now use your cluster with:
```

```
kubectl cluster-info --context kind-tcpdump-cluster
```

Have a nice day! 🎉

```
$ kind get clusters
```

```
tcpdump-cluster
```

Note: remember that for **kubectl cluster-info** command, we need to prepend our cluster name with **kind-**

```
$ kubectl cluster-info --context kind-tcpdump-cluster
```

Similarly unless we used the default cluster name “kind” we need to pass in **--name** for invocations to the kind cli.

```
$ kind get kubeconfig
```

```
ERROR: could not locate any control plane nodes
```

```
$ kind get kubeconfig --name tcpdump-cluster
```

```
apiVersion: v1  
clusters:  
. . .
```

Build tcpdump image

We'll create a container image for our tcpdump image. This can be useful in many circumstances when there are issues and we want to have a container loaded into our cluster for troubleshooting.

In this experiment, we'll build the image with a **"here"** document, similar to the way we built the cluster in a previous experiment.

Paste the following into our terminal

```
docker build -t tcpdump - <<EOF  
FROM ubuntu  
RUN apt-get update && apt-get install -y tcpdump  
CMD tcpdump -i eth0  
EOF  
  
$ docker build -t tcpdump - <<EOF  
> FROM ubuntu  
> RUN apt-get update && apt-get install -y tcpdump  
> CMD tcpdump -i eth0  
> EOF  
#2 [internal] load .dockerignore  
#2  
sha256:6d12e025fda653c659974c1f3ffa6c64eeec150e37e3f3d45512a7a7b  
45bdfc9  
#2 transferring context: 2B 0.0s done  
#2 DONE 0.1s  
  
#1 [internal] load build definition from Dockerfile  
#1  
sha256:a3d4f9355937a661add8618395d3740cfc88354763fdb1424719c309f  
f3564e7  
#1 transferring dockerfile: 121B 0.0s done  
#1 DONE 0.1s  
.  
.  
.  
#6 naming to docker.io/library/tcpdump done
```

#6 DONE 0.2s

Use 'docker scan' to run Snyk tests against images to find vulnerabilities and learn how to fix them

Verify the image

\$ docker images

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
tcpdump	latest	57d9ee3ae218	2 hours ago	109MB
kindest/node	<none>	32b8b755dee8	2 months ago	1.12GB

Since this is a latest we would need to tag that or have to set image pull policy. If we post a latest to the kind cluster it would not behave as we expect due to the default pull policy for latest to always pull.

Tag the image with a unique valid tag.

\$ docker tag tcpdump:latest tcpdump:new

Validate our images

\$ docker images

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
tcpdump	latest	57d9ee3ae218	2 hours ago	109MB
tcpdump	new	57d9ee3ae218	2 hours ago	109MB
kindest/node	<none>	32b8b755dee8	2 months ago	1.12GB

Load the image to our cluster.

kubernetes@DESKTOP-1M2VN7E MINGW64 /c/projects/kind

\$ kind load docker-image tcpdump:new

Image: "tcpdump:new" with ID
"sha256:57d9ee3ae218cd2a637a83a4ad8b5a1de208d52d933358e60c60ca2f
fa4950d7" not yet present on node "kind-control-plane",
loading...

View the images on the cluster with the **crictl** tool

kubernetes@DESKTOP-1M2VN7E MINGW64 /c/projects/kind

\$ docker exec -it kind-control-plane crictl images

Note: If we receive an error we can prepend the command with winpty on windows environments for git bash.

```
kubernetes@DESKTOP-1M2VN7E MINGW64 /c/projects/kind
```

```
$ winpty docker exec -it kind-control-plane crictl images
```

IMAGE	TAG
IMAGE ID	
SIZE	
docker.io/kindest/kindnetd 6de166512aa22	v20210326-1e038dc5
54MB	
docker.io/library/tcpdump 57d9ee3ae218c	new
112MB	
docker.io/rancher/local-path-provisioner e422121c9c5f9	v0.0.14
13.4MB	
k8s.gcr.io/build-image/debian-base c7c6c86897b63	v2.1.0
21.1MB	
k8s.gcr.io/coredns/coredns 296a6d5035e2d	v1.8.0
12.9MB	
k8s.gcr.io/etcd 0369cf4303ffd	3.4.13-0
86.7MB	
k8s.gcr.io/kube-apiserver 94ffe308aef9	v1.21.1
127MB	
k8s.gcr.io/kube-controller-manager 96a295389d472	v1.21.1
121MB	
k8s.gcr.io/kube-proxy 0e124fb3c695b	v1.21.1
133MB	
k8s.gcr.io/kube-scheduler 1248d2d503d37	v1.21.1
51.9MB	
k8s.gcr.io/pause ed210e3e4a5ba	3.5
301kB	

Note: We will also see the kubernetes control plane images, the kindest node, and the image we loaded.

This would enable us to execute the tcpdump to link output the communication occurring between containers.

Experiment Cleanup

Delete the cluster

\$ kind delete cluster --name tcpdump-cluster

Deleting cluster "tcpdump-cluster" ...